



UNIVERSITÀ  
DEGLI STUDI  
DI TRIESTE



# NEW APPROACHES IN SCIENTIFIC COMPUTING: THE DESIGN OF A CLOUD-READY HPC INFRASTRUCTURE

---

*Isac Pasianotto*

Supervised by: dr. Ruggero Lot  
Cosupervised by: prof. Stefano Cozzini

Graduation session: M.Sc. in *Data Science and Scientific Computing*

20th September 2024

## 🤔 What is this thesis about?

*Investigating the possibility to adopt a **cloud-native** approach based on container orchestration in a private **HPC cluster** to provide it in a **SaaS** fashion*

“Cloud computing has been coined as an umbrella term to describe a category of sophisticated **on-demand computing services** initially offered by commercial providers”

– R. Buyya et al. *Cloud Computing: Principles and Paradigms*



# What is “Cloud Computing”



“Cloud computing has been coined as an umbrella term to describe a category of sophisticated **on-demand computing services** initially offered by commercial providers”

– R. Buyya et al. *Cloud Computing: Principles and Paradigms*



for us:

Cloud Computing is the result of the **evolution** of the computing concept **driven by the technology** improvements and by users' requirements.

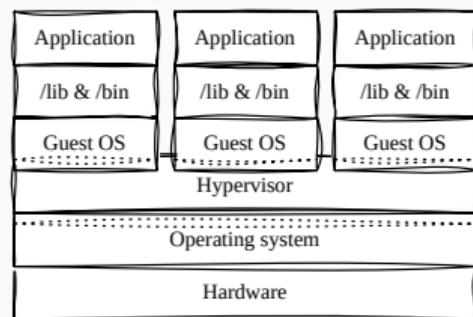


## Virtualization

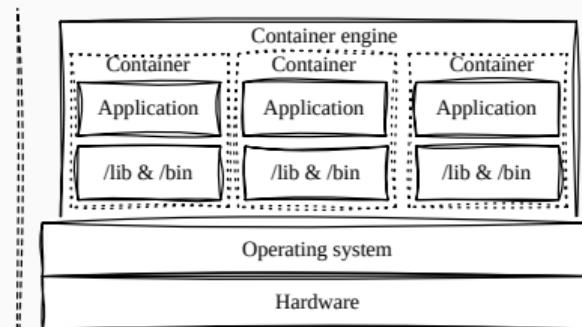
- Isolated environment with its own OS and resources
- **Overhead** for the **hypervisor**

## Containerization

- Host's **processes** → they share the Kernel with the host
- Build over **namespaces**, **cgroups**, and **capabilities**



Virtualization



Containerization

## Performance/Usability trade-off

**Container orchestrator:** simplify the interaction with the cluster for both the user and the administrator

- Automatic scaling and load balancing
- Availability, self healing and fault tolerance
- **Standardized API** → portability
- ...



**kubernetes**

## Performance/Usability trade-off

**Container orchestrator:** simplify the interaction with the cluster for both the user and the administrator

- Automatic scaling and load balancing
- Availability, self healing and fault tolerance
- **Standardized API** → portability
- ...



**kubernetes**



*First attempts in this direction were catastrophic*



*Network latency rapidly became the bottleneck of the system*

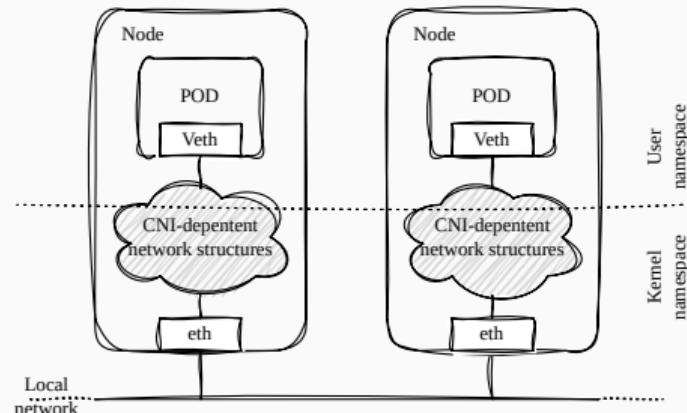
 Def.

A CNI plugin is responsible for **inserting a network interface** into the **container network namespace** (e.g., one end of a virtual ethernet (veth) pair) and making any **necessary changes** on the host (e.g., attaching the other end of the veth into a bridge).

– Red Hat sysadmin's guide



*“And making any necessary changes on the host . . .”*





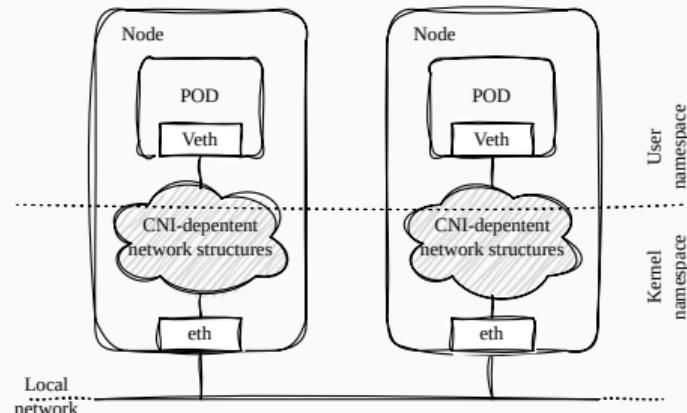
*“And making any necessary changes on the host . . .”*



AREA  
SCIENCE PARK



→ **Calico: BGP routing**



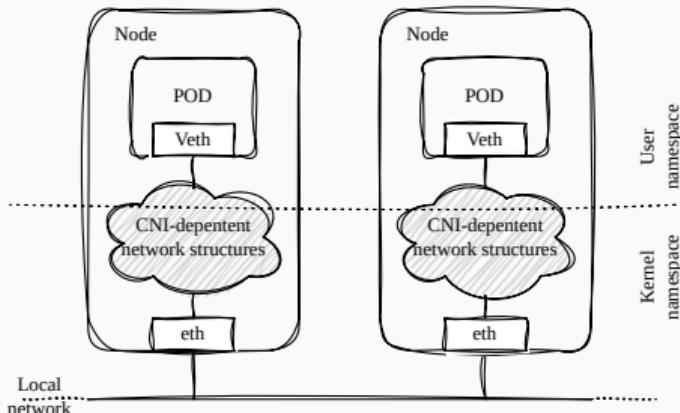


*“And making any necessary changes on the host . . .”*



Calico: BGP routing

→ Flannel: VXLAN encapsulation





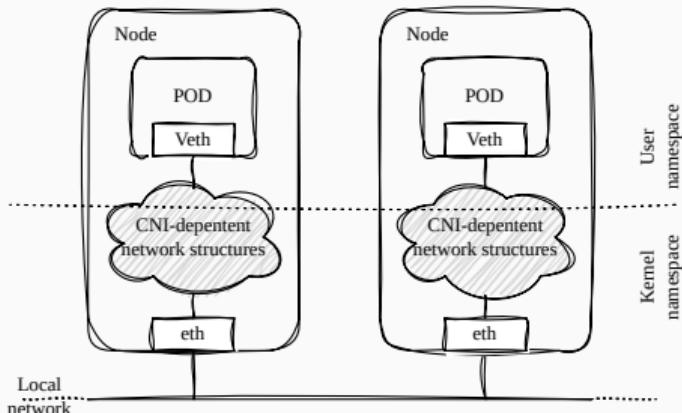
*“And making any necessary changes on the host . . .”*



*Calico:* **BGP** routing

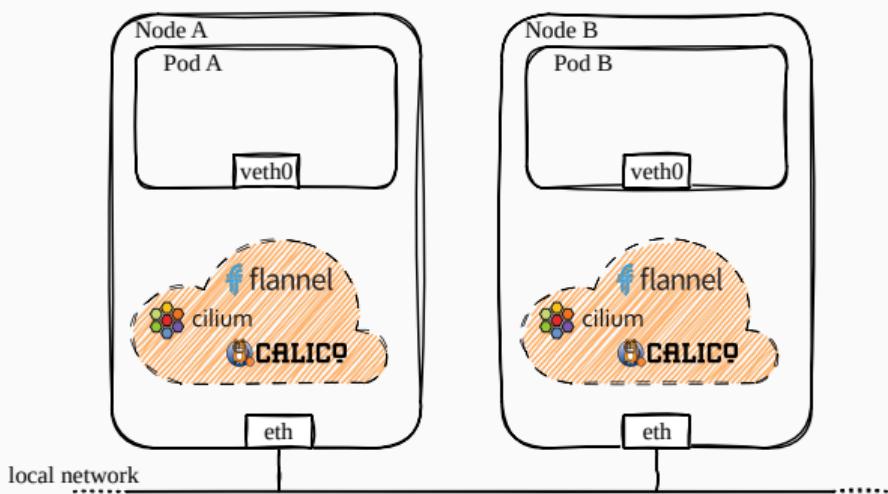
*Flannel:* **VXLAN** encapsulation

→ *Cilium:* **eBPF** (sandbox technology within the Linux kernel)



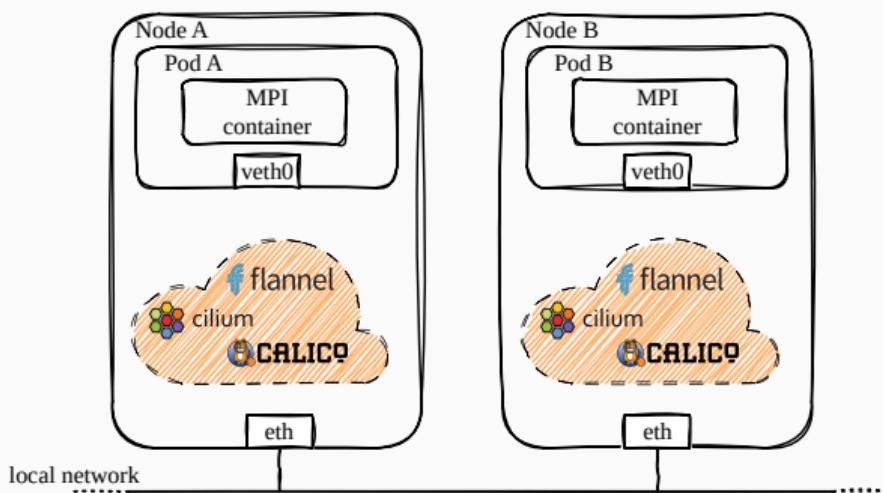
## The idea:

- Spawn 2 pods in 2 (different) nodes



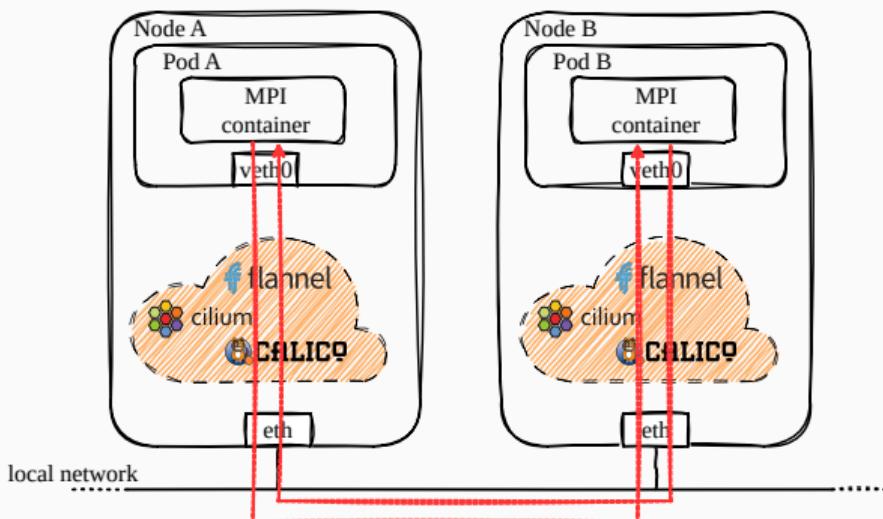
## The idea:

- Spawn 2 pods in 2 (different) nodes
- Each pod will have a container running an MPI process



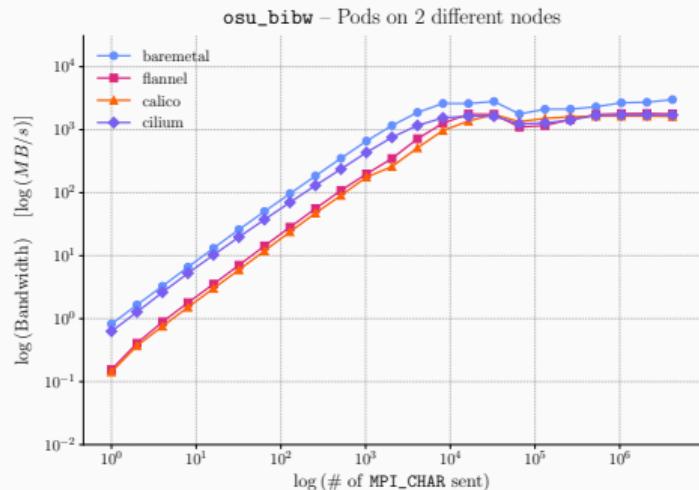
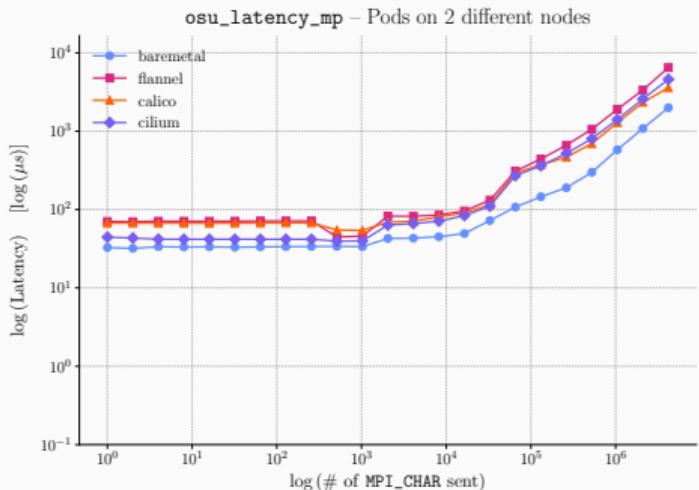
## The idea:

- Spawn 2 pods in 2 (different) nodes
- Each pod will have a container running an MPI process
- Perform ping-pong tests between the 2 pods 🎾



# Results

🔍 Nodes with 25 Gbps ethernet ports



	Bare metal	Calico	Flannel	Cilium
Min	32.7	67.5	70.4	44.4

	Bare metal	Calico	Flannel	Cilium
Max	3.016e3	1.62e3	1.78e3	1.71e3

## 💡 What is Dask?

- Flexible library for **parallel computing** in **Python**
- Allow to *distribute the computation without managing the parallelism*

## 💡 What is Dask?

- Flexible library for **parallel computing** in **Python**
- Allow to *distribute the computation without managing the parallelism*

### 💡 What is Dask?

- Flexible library for **parallel computing** in **Python**
- Allow to *distribute the computation* without managing the parallelism

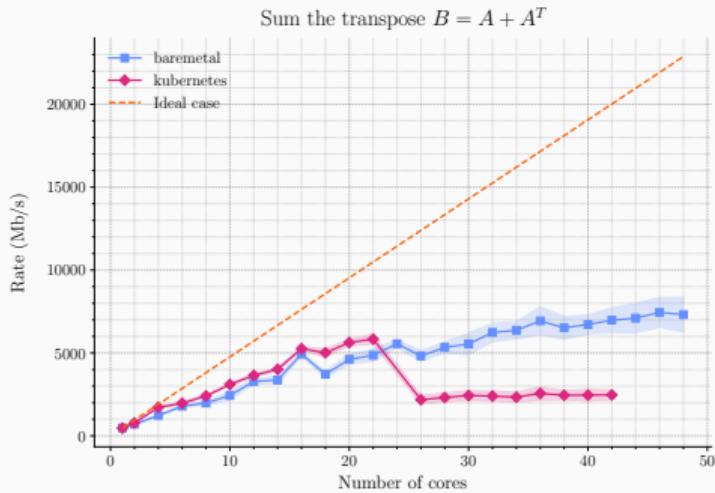
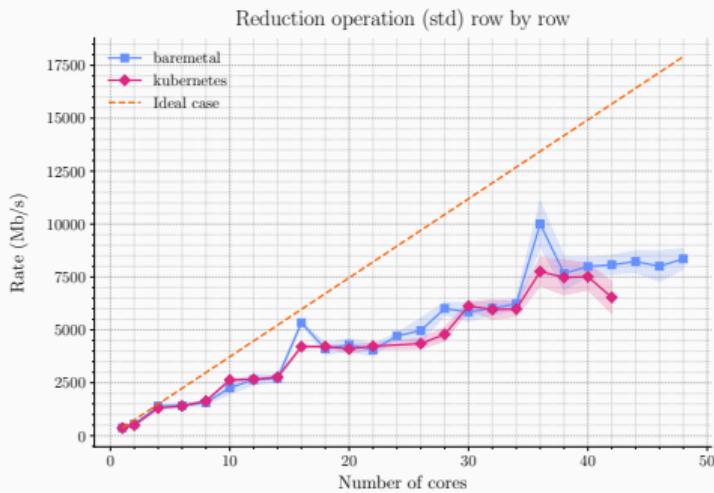
💡 Our solution to develop a parallel code to solve **most common problems in data science!**



# The results



🔍 Nodes with  $2 \times 12$  cores each.





To sum up:



AREA  
SCIENCE PARK



🤔 Can we use Kubernetes to deploy and manage HPC applications?



## To sum up:



🤔 Can we use Kubernetes to deploy and manage HPC applications?

- 🎉 **Yes!**, but with *some limitation*.



## To sum up:



🤔 Can we use Kubernetes to deploy and manage HPC applications?

- 🎉 **Yes!**, but with *some limitation*.
- Times are **not** mature yet for a **full migration**
- ⚠ In the **communication bounded** application, the performance gap is **not acceptable**
- In all the other cases, the **overhead** is **negligible**



## To sum up:



🤔 Can we use Kubernetes to deploy and manage HPC applications?

- 🎉 **Yes!**, but with *some limitation*.
- Times are **not** mature yet for a **full migration**
- ⚠️ In the **communication bounded** application, the performance gap is **not acceptable**
- In all the other cases, the **overhead** is **negligible**
- Kubernetes** comes with a lot of benefits that improve the user experience
- 👉 The technology evolution of last years make us be **optimistic** about the future

*Thank you for your attention*